

Stanford CS 224n Assignment 2

Hanchung Lee

January 19, 2020

1 Written: Understanding word2Vec (23 points)

(a) (3 points) Show that the naive-softmax loss given in Equation (2) is the same as the cross-entropy loss between y and \hat{y} ; i.e., show that

$$-\sum_{w \in Vocab} y_w \log \hat{y}_w = -\log \hat{y}_o$$

Your answer should be one line.

Answer: Since y is a one-hot vector where $y_w = 1$ when $w = o$ and $y_w = 0$ when $w \neq o$,

$$\begin{aligned} -\sum_{w \in Vocab} y_w \log(\hat{y}_w) &= -[y_1 \log(\hat{y}_1) + \dots + y_o \log(\hat{y}_o) + \dots + y_w \log(\hat{y}_w)] \\ &= -y_o \log(\hat{y}_o) \\ &= -\log(\hat{y}_o) \end{aligned}$$

(b) (5 points) Compute the partial derivative of $J_{naive-softmax}(v_c, o, U)$ with respect to v_c . Please write your answer in terms of y , \hat{y} , and U .

Answer: Let input vector be $\theta = U^\top v_c$ and prediction function be $\hat{y} = softmax(\theta)$. From (a), we know the derivative of cross-entropy loss is equivalent to softmax loss for one hot vector \mathbf{y} , therefore,

$$\begin{aligned} J &= CrossEntropy(y, \hat{y}) \\ \therefore \frac{\partial J}{\partial \theta} &= (\hat{y} - y)^\top \end{aligned}$$

Reference: <https://deeppnotes.io/softmax-crossentropy>

From above, we can use chain rule to solve the derivative:

$$\begin{aligned} \frac{\partial J}{\partial v_c} &= \frac{\partial J}{\partial \theta} \frac{\partial \theta}{\partial v_c} \\ &= (\hat{y} - y)^\top \frac{\partial U^\top v_c}{\partial v_c} \\ &= U(\hat{y} - y) \end{aligned}$$

(c) (5 points) Compute the partial derivatives of $J_{naive-softmax}(\mathbf{v}_c, o, \mathbf{U})$ with respect to each of the ‘outside’ word vectors, u_w ’s. There will be two cases: when $w = o$, the true ‘outside’ word vector, and $w \neq o$, for all other words. Please write your answer in terms of y , \hat{y} , and v_c .

Answer: Similar to answer (b) above.

$$\begin{aligned}\frac{\partial J}{\partial v_c} &= \frac{\partial J}{\partial \theta} \frac{\partial \theta}{\partial v_c} \\ &= (\hat{y} - y) \frac{\partial \mathbf{U}^\top v_c}{\partial \mathbf{U}} \\ &= (\hat{y} - y) v_c\end{aligned}$$

(d) (3 Points) The sigmoid function is given by Equation 4:

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1} \quad (1)$$

Please compute the derivative of $\sigma(x)$ with respect to x , where x is a scalar. Hint: you may want to write your answer in terms of $\sigma(x)$.

Answer: Apply quotient rule $f(x) = \frac{g(x)}{h(x)}$, $f'(x) = \frac{g'(x)h(x) + g(x)h'(x)}{h^2(x)}$, we have:

$$\begin{aligned}\frac{\partial \sigma(x)}{\partial x} &= \frac{\partial \frac{e^x}{e^x + 1}}{\partial x} \\ &= \frac{(e^x + 1)e^x - e^x e^x}{(e^x + 1)^2} \\ &= \frac{e^x}{(e^x + 1)^2} \\ &= \sigma(x) \frac{1}{e^x + 1} \\ &= \sigma(x) \frac{-e^x + e^x + 1}{e^x + 1} \\ &= \sigma(x) \left(\frac{-e^x}{e^x + 1} + \frac{e^x + 1}{e^x + 1} \right) \\ &= \sigma(x)(1 - \sigma(x))\end{aligned}$$

(e) (4 points) Now we shall consider the Negative Sampling loss, which is an alternative to the Naive Softmax loss. Assume that K negative samples (words) are drawn from the vocabulary. For simplicity of notation we shall refer to them as w_1, w_2, \dots, w_K and their outside vectors as $\mathbf{u}_1, \dots, \mathbf{u}_K$. Note that $o \notin w_1, \dots, w_K$. For a center word c and an outside word o , the negative sampling loss function is given by:

$$J_{neg-sample}(v_c, o, U) = -\log(\sigma(u_o^\top v_c)) - \sum_{k=1}^K \log \sigma(-u_k^\top v_c)$$

for a sample w_1, \dots, w_K , where $\sigma(\cdot)$ is the sigmoid function.

Please repeat parts (b) and (c), computing the partial derivatives of $J_{neg-sample}$ with respect to v_c , with respect to u_o , and with respect to a negative sample u_k . Please write your answers in terms of the vectors u_o , v_c , and u_k , where $k \in [1, K]$. After you've done this, describe with one sentence why this loss function is much more efficient to compute than the naive-softmax loss. Note, you should be able to use your solution to part (d) to help compute the necessary gradients here.

Answer:

$$\begin{aligned}\frac{\partial J}{\partial u_o} &= (\sigma(u_o^\top v_c) - 1)v_c \\ \frac{\partial J}{\partial u_k} &= (1 - \sigma(-u_k^\top v_c))v_c, \forall k \in [1, K] \\ \frac{\partial J}{\partial v_c} &= (\sigma(u_o^\top v_c) - 1)u_o - \sum_{k=1}^K (\sigma(-u_k^\top v_c) - 1)u_k\end{aligned}$$

Answer: For naive softmax loss, it computes the whole outside vectors U . But for negative sampling loss, it only calculates a fixed size K . Thus, negative sampling loss is more compute and memory efficient

(f) (3 points) Suppose the center word w_t and the context window is $[w_{t-m}, \dots, w_{t-1}, w_t, w_{t+1}, \dots, w_{t+m}]$, where m is the context window size. Recall that for the skip-gram version of word2Vec, the total loss for the context window is:

$$J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}, U) = \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} J(v_c, w_{t+j}, U) \quad (2)$$

Here, $J(v_c, w_{t+j}, U)$ represents an arbitrary loss term for the center word $c = w_t$ and outside word w_{t+j} . $J(v_c, w_{t+j}, U)$ could be $J_{\text{naive-softmax}}(v_c, w_{t+j}, U)$ or $J_{\text{neg-sample}}(v_c, w_{t+j}, U)$, depending on your implementation.

Write down three partial derivatives:

- (i) $\partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}) / \partial U$
- (ii) $\partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}) / \partial v_c$
- (iii) $\partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}) / \partial w_c$ when $w \neq c$

Write your answers in terms of $\partial J(v_c, w_{t+j}, U) / \partial U$ and $\partial J(v_c, w_{t+j}, U) / \partial v_c$. This is very simple – each solution should be one line.

Answer:

$$\begin{aligned} \partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}) / \partial U &= \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \frac{J(v_c, w_{t+j}, U)}{\partial U} \\ \partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}) / \partial v_c &= \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \frac{J(v_c, w_{t+j}, U)}{\partial v_c} \\ \partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}) / \partial w_c (w \neq c) &= 0 \end{aligned}$$